

# Intro to Data Analysis with R

Pri Oberoi

5/16/2016

Make sure you have all the content in this github repo:  
[www.github.com/prioberoi/R\\_intro\\_to\\_data\\_analysis](https://www.github.com/prioberoi/R_intro_to_data_analysis)




**Pri Oberoi** ([poberoi@doc.gov](mailto:poberoi@doc.gov))  
Data Scientist, Commerce Data Service  
US Department of Commerce

# Goals

Walk away with the foundations for

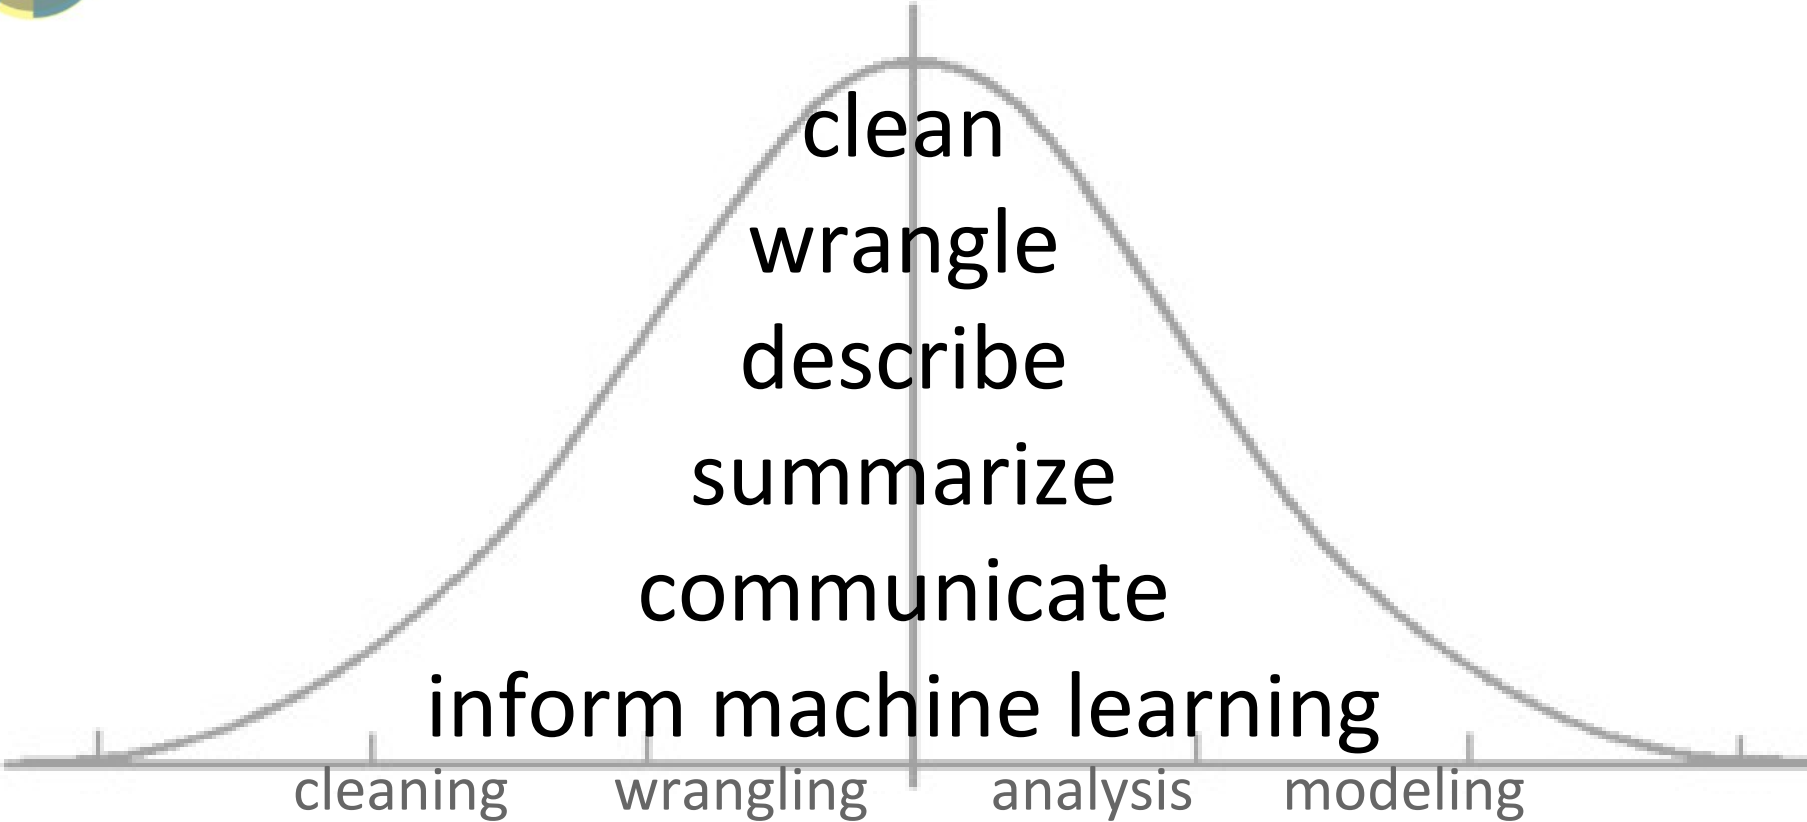
- The role of data analysis is in the data science pipeline
- R markdown
- Data visualizations
- Clean data
- Aggregate and summarize data
- Statistical tests

# The Data Science Pipeline



# Why do data analysis?

clean  
wrangle  
describe  
summarize  
communicate  
inform machine learning



clean  
wrangle  
describe  
summarize  
communicate  
inform machine learning


cleaning

wrangling

analysis

modeling

# The Pipeline





# R Markdown

## Output formats:

HTML, PDF, MS Word, HTML5 slides,  
books, dashboards, websites

## Benefits:

Easy to create

Embedded R code chunks

(which can be visible or not on the final output)

Allows you to add a narrative through your code

Reproducible

RStudio

Project: (None)

data\_analysis\_with\_R.Rmd \* Untitled1\* \*

Go to file/function Addins

Knit HTML

```
1 ---
2 title: "Intro to Data Analysis with R"
3 author: "Pri Oberoi"
4 date: "May 16, 2016"
5 output: html_document
6 ---
7
8 # Intro to Data Analysis with R
9
10 ```{r, echo=TRUE, message=FALSE}
11
12 packagesNeeded <- c("ggplot2", "reshape2", "Rmisc")
13 packagesToInstall <- packagesNeeded[!(packagesNeeded %in% installed.packages[, "Package"]) ]
14 if(length(packagesToInstall)) install.packages(packagesToInstall)
15
16 library(ggplot2)
17 library(reshape2)
18 library(Rmisc)
19 ```
20
21 ## Visualization
22
23 qplot() is good for quick plots and is similar to plot()
24 ggplot() is more verbose but it has more functionality
25
26 ### Scatter plots
27 ```{r, echo = TRUE}
16:17 [ Chunk 1 ] R Markdown
```

Code chunk

Run chunk

Navigate between chunks

Environment History

Global Environment

Data

|              |                               |
|--------------|-------------------------------|
| data         | 149808 obs. of 17 variables   |
| data_clean   | 599232 obs. of 7 variables    |
| dd           | chr [1, 1:205] "Thu May 12... |
| df           | chr [1, 1:50] "Thu May 12 ... |
| fd           | chr [1, 1:205] "Thu May 12... |
| residuals... | 149808 obs. of 2 variables    |

Values

|                   |                                 |
|-------------------|---------------------------------|
| c                 | Classes 'url', 'connection' ... |
| corTest           | List of 9                       |
| fit               | Large lm (12 elements, 32 Mb)   |
| fit\$resid        | Large numeric (149808 elemen... |
| i                 | 2L                              |
| index             | 3L                              |
| j                 | Large list (4634 elements, 1... |
| :List of 25       |                                 |
| ..\$ created_at : | chr "Thu May 12 18:52:...       |

Files Plots Packages Help Viewer

R: Data frame

Find in Topic

In versions of R prior to 2.4.0 row.names had to be character to ensure compatibility with such versions of R, supply a character vector as the row.names argument.

References

Chambers, J. M. (1992) *Data for models*. Chapter 3 of *Statistical Models in S* eds J. M. Chambers and T. J. Hastie, Wadsworth & Brooks/Cole.

See Also

Console R Markdown \*

```
[952] "quoted_status.coordinates.coordinates2"
[953] "retweeted_status.quoted_status.geo.type"
[954] "retweeted_status.quoted_status.geo.coordinates1"
[955] "retweeted_status.quoted_status.geo.coordinates2"
[956] "retweeted_status.quoted_status.coordinates.type"
[957] "retweeted_status.quoted_status.coordinates.coordinates1"
[958] "retweeted_status.quoted_status.coordinates.coordinates2"
> |
```

Console output appears here

Dataframes and other objects appear here

Plots appear here

# Analysis Toolkit

Visualization

Statistics

Aggregation

# Data Visualization

# ggplot2

qplot()

“quick plot”

- similar to plot() from base R
- less typing
- less customizable

```
qplot(carat, price, data = diamonds,  
size = I(1), alpha = I(1/10), main =  
"qplot scatter plot")
```

ggplot()

- more customizable
- more functionality

```
ggplot(data = diamonds, aes(x = carat,  
y = price)) +  
  geom_point(size = 1, alpha = 1/10) +  
  ggtitle("ggplot scatter plot")
```

# ggplot2

x

y

```
qplot(carat, price, data = diamonds,  
size = I(1), alpha = I(1/10), main =  
"qplot scatter plot")
```


aesthetics like  
point size, point  
transparency,  
plot title

data

y

x

```
ggplot(data = diamonds, aes(x = carat,  
y = price)) +  
geom_point(size = 1, alpha = 1/10) +  
ggtitle("ggplot scatter plot")
```



## Your turn (10 mins)

Run the code in chunk 2: **Scatterplots**

Update `ggplot()` code so the color of the scatterplot points varies based on the value of 'cut'

You can do this by adding a 'colour =' argument to the `aes()` mapping



## Your turn (10 mins)

Run the code in chunk 3: **Histograms and Bar Charts**

Note that you can set the 'binwidth' for histograms

Run the code in chunk 4: **Boxplots and Violin Plots**

Box plots: more widely interpretable

Violin plots: useful for non-normal distributions and to scale to  
number of observations

# NTIA Broadband Data Example

NTIA's broadband data from June, 2014 for Washington, DC

# Hypothesis Testing

Null hypothesis: the typical upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

# Your turn (10 mins)

Import the data by running chunk 5

Look at the dataframe

```
View(data)
dim(data)
names(data)
str(data)
summary(data)
```

Create a histogram of max advertised download speeds  
(maxaddown)

10 min break

# Cleaning

# Messy Data

Signs you have messy data:

- Column headers are values, not variable names
- Multiple variables are stored in one column
- Variables are stored in both rows and columns
- Multiple types of observational units are stored in the same table
- A single observational unit is stored in multiple tables



[This content is from Hadley Wickham's paper on tidy data](#)

Column headers are values, not variable names

We will be looking at the maxaddown, maxadup, typicdown, typicup variables

They are stored as different columns/variables, rather than different values.

```
# Run chunk 6
```

```
melt() # this is a function that converts columns into rows  
?melt # use ? to read the documentation on a function
```



Multiple variables are stored in one column

`data_clean` now has one column named 'variable' that contains the variable indicating if this is advertised or typical as well as whether this speed is for uploads or downloads.

```
# Run chunk 7 and look at the resulting data_clean dataframe
```

```
data_clean$speedDirection <- "download"  
data_clean$speedDirection[data_clean$variable %in% c("maxadup","typicup")]  
<- "upload"  
data_clean$speedSource <- "advertised"  
data_clean$speedSource[data_clean$variable %in% c("typicup","typicdown")]  
<- "typical"  
data_clean$variable <- NULL
```

Variables are stored in both rows and columns

We don't have this problem in our dataset, but here is an example:

| id      | year | month | element | d1 | d2   | d3   | d4 | d5   | d6 | d7 | d8 |
|---------|------|-------|---------|----|------|------|----|------|----|----|----|
| MX17004 | 2010 | 1     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 1     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 2     | tmax    | —  | 27.3 | 24.1 | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 2     | tmin    | —  | 14.4 | 14.4 | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 3     | tmax    | —  | —    | —    | —  | 32.1 | —  | —  | —  |
| MX17004 | 2010 | 3     | tmin    | —  | —    | —    | —  | 14.2 | —  | —  | —  |
| MX17004 | 2010 | 4     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 4     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 5     | tmax    | —  | —    | —    | —  | —    | —  | —  | —  |
| MX17004 | 2010 | 5     | tmin    | —  | —    | —    | —  | —    | —  | —  | —  |

Table 11: Original weather dataset. There is a column for each possible day in the month. Columns d9 to d31 have been omitted to conserve space.

Variables are individual columns (id, year, month), spread across columns (day, d1–d31) and across rows (tmin, tmax)

Multiple types of observational units are stored in the same table

data\_clean has data on different observational units, the provider/holding company, broadband speeds, location

Repeating values in a column are a result.

Similar to database normalization.

Note, some data analysis tools work with denormalized data

A single observational unit is stored in multiple tables


Data values for a single variable are found across tables

```
# the following functions are helpful  
rbind() #add dataframe as rows, must have same number of columns  
cbind() #add dataframe as columns, must have same number of rows  
merge() #merge two data frames using an identifier column  
ldply() #from the plyr package reads multiple csvs into one dataframe
```

# Summarizing Data Within Groups

# Summarizing Data Within Groups

This boxplot indicates advertised and typical speeds differ.



```
# aggregate() will run functions over a group you specify
```

```
# what does this do:
```

```
# data_clean[data_clean$speedDirection == 'download','value']
```

```
aggregate(data_clean[data_clean$speedDirection == 'download','value'], list  
(data_clean[data_clean$speedDirection == 'download', 'speedSource']), mean)
```

## Your turn (5 mins)

Update chunk 8 to show the mean upload speeds by  
speedSource (advertised or typical)

10 min break




# Correlation

# Correlation

We expect that the advertised and typical speeds are correlated


Run chunk 9

# Correlation



Pearson's product-moment correlation

```
data: data$maxaddown and data$typicdown  
t = 242.74, df = 149810, p-value < 2.2e-16  
alternative hypothesis: true correlation is not equal to 0  
95 percent confidence interval:  
 0.5276687 0.5349375  
sample estimates:  
      cor  
0.5313129
```




p-value < 0.05  
not surprising

# Revisit our null hypothesis

Null hypothesis: the typical upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

# Comparing Samples



Null hypothesis: the typical download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Let's do a quick t-test to see if that difference is statistically significant

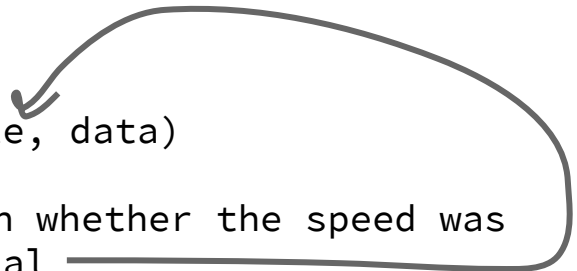
```
t.test(response ~ variable, data)
```

Null hypothesis: the typical download speeds for broadband providers in Washington, DC are the same as the advertised speeds

Let's do a quick t-test to see if that difference is statistically significant



```
t.test(response ~ variable, data)
```



#does the download speed differ based on whether the speed was advertised or typical in our dataset?

Null hypothesis: the typical mean upload and download speeds for broadband providers in Washington, DC are the same as the advertised speeds

$H_0$

## Run chunk 10

```
> t.test(data_clean[data_clean$speedDirection == 'download', 'value'] ~ data_clean[data_clean$speedDirection == 'download', 'speedSource'], data_clean)
```

Welch Two Sample t-test

```
data: data_clean[data_clean$speedDirection == "download", "value"] by data_clean[data_clean$speedDirection == "download", "speedSource"]
t = 266.22, df = 243900, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 2.031413 2.061546
sample estimates:
mean in group advertised      mean in group typical
      5.346951              3.300471
```

P-value < 0.05

$H_a$



## Your turn! (10 mins)

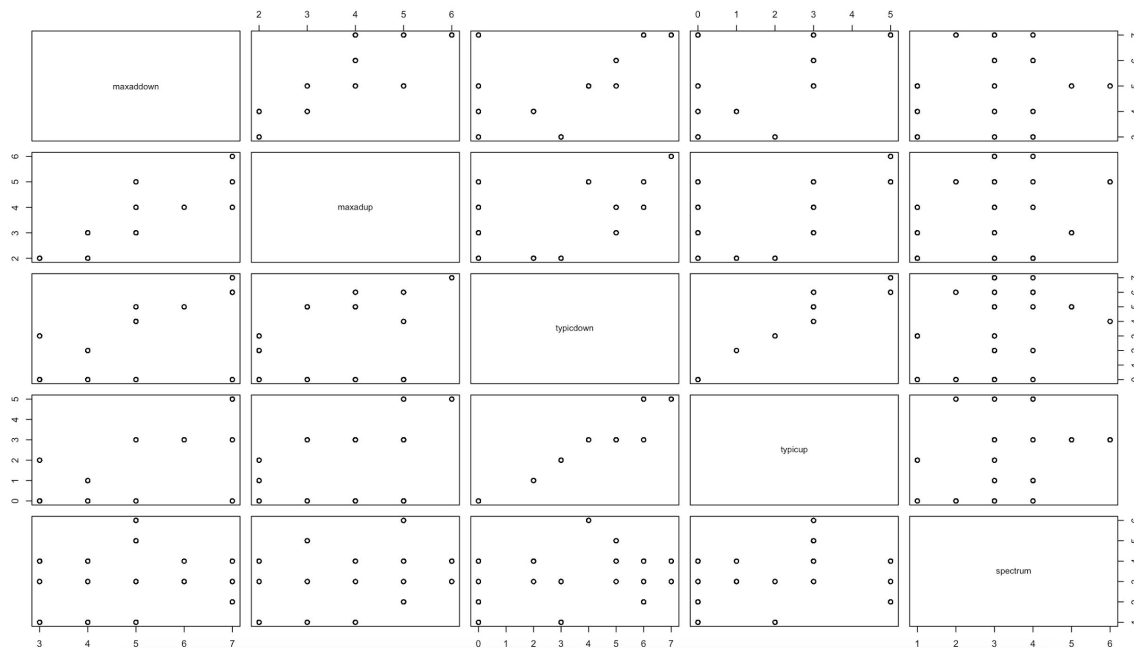
update the t-test code (chunk 10) to see if the typical upload speed is different than the advertised upload speed

10 min break

# Analysis for Model Selection and Feature Selection


# Scatterplot matrices

Relationship between continuous variables in your data




# Linear regression

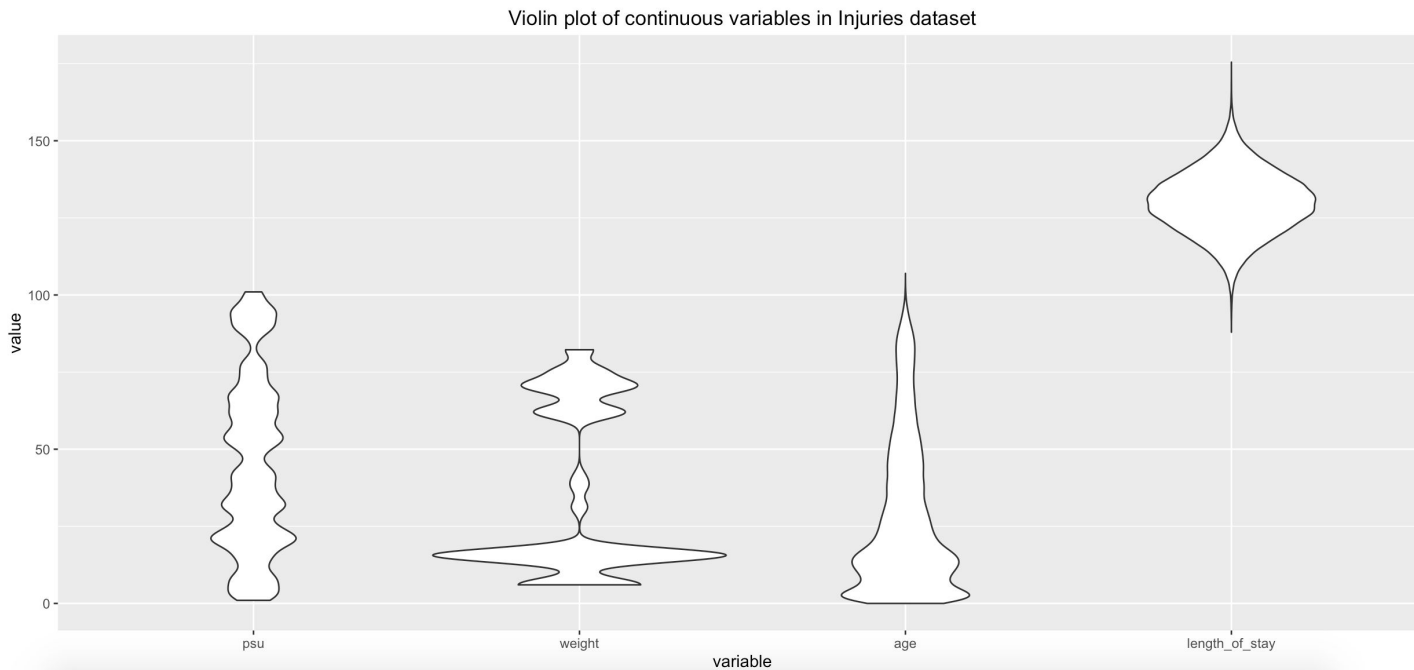
Scatterplot of typicdown and maxaddown



Scatterplot of typicdown and spectrum




# Normalizing data and missing data



# Normalized

Violin plot of noralized continuous variables in Injuries dataset

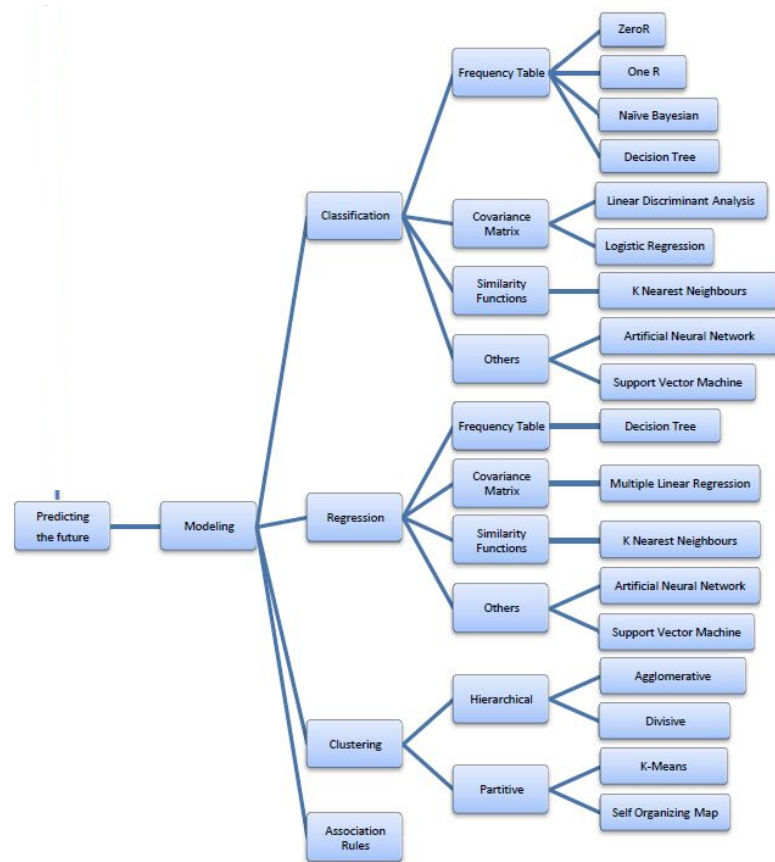
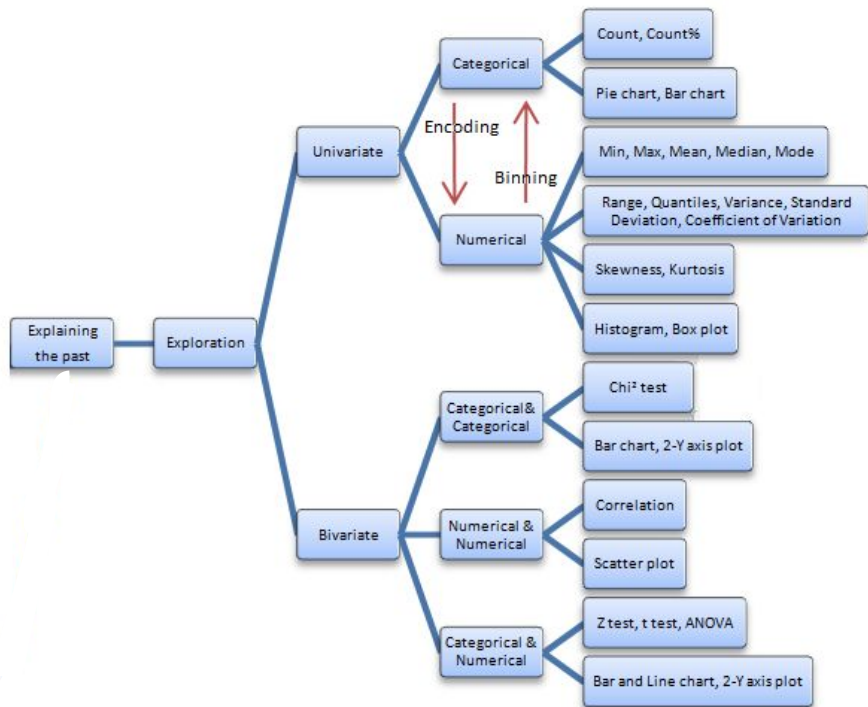


# From statistical tests to statistical learning

Statistical tests describe your data

Statistical learning allows you to make predictions about the future





# Statistical learning allows us to ask questions like:

Can we predict what the typical upload and download speeds are, for a given provider, if we know the advertised upload and download speeds?

## Resources

**R-bloggers:** <http://www.r-bloggers.com/>

**FlowingData:** <http://flowingdata.com/category/tutorials/>

**Google's R Style Guide:** <https://google.github.io/styleguide/Rguide.xml>

**R Markdown:** <http://rmarkdown.rstudio.com/>

**Saed Sayad's data mining map:** [http://www.saedsayad.com/data\\_mining\\_map.htm](http://www.saedsayad.com/data_mining_map.htm)

**See github repo for more resources!**